

Compositional Verification of Stigmergic Collective Systems

Luca Di Stefano¹ Frédéric Lang²

¹University of Gothenburg, Sweden

²INRIA Grenoble, France

VMCAI'2023



CHALMERS
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF GOTHENBURG

Inria

- Collections of agents interacting with each other
- Found in CS, economics, biology. . .
- Interactions and feedback may lead to emergence of collective features
- Reasoning about emergence is hard. Model checking?

- Collections of agents interacting with each other
- Found in CS, economics, biology...
- Interactions and feedback may lead to emergence of collective features
- Reasoning about emergence is hard. Model checking?

Pros

- Can prove emergence, safety, etc. (arbitrary temporal properties)
- Push-button, no human guidance needed

Cons

- Requires user expertise for modelling, specification
- State space explosion as the number/complexity of agents increases

What LAbS is about

- High-level language to concisely specify systems/properties
- Focus on indirect, attribute-based interaction mechanisms
- Reuse of different existing verification technologies
 - E.g., CADP, which offers model-checking and compositional verification tools out of the box

What LAbS is about

- High-level language to concisely specify systems/properties
- Focus on indirect, attribute-based interaction mechanisms
- Reuse of different existing verification technologies
 - E.g., CADP, which offers model-checking and compositional verification tools out of the box

Contributions


- Encoding of LAbS systems into parallel LNT programs
- Compositional verification workflow
- Sound value analysis to prune individual state spaces and speed up verification

```
1 system {
2   spawn = Node: N
3 }
4 stigmergy Election {
5   link = true
6   leader: N
7 }


8 agent Node {
9   stigmergies = Election
10  Behavior =
11    leader >= id ->
12      leader <~ id;
13    Behavior
14 }
```

- N nodes run for election, by storing their `id` in a *stigmergic* variable `leader`. If `leader < id`, the node waits
- All communication is implicit
 - Nodes exchange their values of `leader`
 - Values are timestamped, “newer” ones replace “older”
 - `link = true` means broadcast messages
- Intuitively, they should converge to a state where all nodes set `leader` to the lowest `id` in the system

Some feasible executions (with $N = 3$)



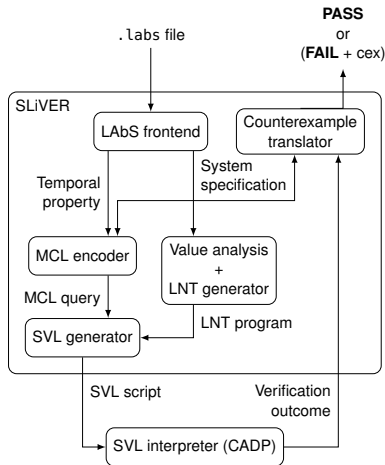
Some feasible executions (with $N = 3$)



A property of interest

$$fairly_{\infty} \forall x : Node \bullet x.leader = 0$$

Along every fair execution, there are infinitely many states where all nodes have 0 as the leader



Based on CADP and its languages/formalisms:

LNT system description

MCL property specification
(alternation-free μ -calculus with data)

SVL scripting of complex verification tasks (in our case: compositional state space generation + model checking)

Program = processes communicating over gates

Send offer

$G(v_1, \dots, v_n)$

Offer values over
gate G

Receive offer

$G(?x_1, \dots, ?x_n)$ where $\varphi(x_1, \dots, x_n)$

Receive n values over G and
bind them to x_i , $i = 1, \dots, n$, *but*
only if $\varphi(x_1, \dots, x_n)$ holds.

Program = processes communicating over gates

Send offer

$G(v_1, \dots, v_n)$

Offer values over
gate G

Receive offer

$G(?x_1, \dots, ?x_n)$ where $\varphi(x_1, \dots, x_n)$

Receive n values over G and
bind them to $x_i, i = 1, \dots, n$, *but*
only if $\varphi(x_1, \dots, x_n)$ holds.

Parallel composition

$\text{par } G_{11}, \dots, G_{1j} \rightarrow P_1 \parallel \dots \parallel G_{n1}, \dots, G_{nk} \rightarrow P_n \text{ end par}$

All processes with G in their set of gates must synchronize over it

Program = processes communicating over gates

Send offer

$G(v_1, \dots, v_n)$

Offer values over
gate G

Receive offer

$G(?x_1, \dots, ?x_n)$ where $\varphi(x_1, \dots, x_n)$

Receive n values over G and
bind them to x_i , $i = 1, \dots, n$, *but*
only if $\varphi(x_1, \dots, x_n)$ holds.

Parallel composition


$\text{par } G_{11}, \dots, G_{1j} \rightarrow P_1 \parallel \dots \parallel G_{n1}, \dots, G_{nk} \rightarrow P_n \text{ end par}$

All processes with G in their set of gates must synchronize over it


Etc.

Loops, nondeterministic choice, conditionals, guards, ...

Each agent is encoded as a process with this structure:



Each agent is encoded as a process with this structure:



- System = Parallel composition of all agents and additional processes (e.g., information about timestamps)
- Multi-party synchronization to resolve the agents' choices

Given a tree of parallel processes S , generate the transition system $lts(S)$ by composing the (minimized) TSs of the “leaves” $P_1, \dots, P_m \in S$

Given a tree of parallel processes S , generate the transition system $Its(S)$ by composing the (minimized) TSs of the “leaves” $P_1, \dots, P_m \in S$

Root-leaf reduction (modulo R)

- For every P_i generate $T_i = Its(P_i)$
- Minimize every T_i modulo R : $T'_i = min_R(T_i)$
- Generate $T = Its(S[T'_i / P_i])$
- Return $min_R(T)$

Given a tree of parallel processes S , generate the transition system $Its(S)$ by composing the (minimized) TSs of the “leaves” $P_1, \dots, P_m \in S$

Root-leaf reduction (modulo R)

- For every P_i generate $T_i = Its(P_i)$
- Minimize every T_i modulo R : $T'_i = min_R(T_i)$
- Generate $T = Its(S[T'_i / P_i])$
- Return $min_R(T)$

Drawback

When generating each T_i we do not know what messages we may receive from the other processes

E.g., in the bully election, *leader* may be any integer in $-128..127$


1. Compute abstract initial state ς_0 from specification \mathbb{S}
 - In this work we use powersets of intervals
2. Add ς_0 to a set σ
3. For every assignment a in \mathbb{S} and every state ς in σ :
 - Evaluate a on ς
 - Add resulting state ς' to σ
4. Reach a fixed point
5. Merge all states in σ to obtain $\bar{\sigma}$

1. Compute abstract initial state ς_0 from specification \mathbb{S}
 - In this work we use powersets of intervals
2. Add ς_0 to a set σ
3. For every assignment a in \mathbb{S} and every state ς in σ :
 - Evaluate a on ς
 - Add resulting state ς' to σ
4. Reach a fixed point
5. Merge all states in σ to obtain $\bar{\sigma}$

Example

In the bully election system we find out that $leader \in \{0, \dots, N\}$ in every state

We use $\bar{\sigma}$ to prune out receptions of impossible messages:



(In practice we plug $\bar{\sigma}$ as a where-clause on receive offers)

System	Baseline		Compositional		Parallel	
	Time (s)	Memory (MiB)	Time (s)	Memory (MiB)	Time (s)	Memory (MiB)
flock-rr	1875	12000	4461	11805	4426	11805
flock	4787	30865	4071	11113	4038	11113
formation-rr	1670	1657	2511	1938	1558	5875
leader5	10	41	34	117	18	212
leader6	77	147	104	225	65	258
leader7	1901	2038	374	404	326	404
twophase2	9	50	67	93	34	210
twophase3	500	209	233	322	131	560

Baseline Previous (sequential) LNT encoding

Compositional Our work spaces on a dedicated core

Parallel What would happen if we split state space generation across multiple cores

-rr Round-robin scheduling

System	Baseline		Compositional		Parallel	
	Time (s)	Memory (MiB)	Time (s)	Memory (MiB)	Time (s)	Memory (MiB)
flock-rr	1875	12000	4461	11805	4426	11805
flock	4787	30865	4071	11113	4038	11113
formation-rr	1670	1657	2511	1938	1558	5875
leader5	10	41	34	117	18	212
leader6	77	147	104	225	65	258
leader7	1901	2038	374	404	326	404
twophase2	9	50	67	93	34	210
twophase3	500	209	233	322	131	560

- *B* wins on very small instances (no overhead)
- *C* scales better and has fewer issues with full interleaving
- *P* brings further gains wrt verification times but may be more memory hungry

Conclusion

- Model checking enables verification of expressive properties in collective systems
- Compositional verification can palliate state space explosion
- Value analysis speeds up state space generation

Conclusion

- Model checking enables verification of expressive properties in collective systems
- Compositional verification can palliate state space explosion
- Value analysis speeds up state space generation

Future work

- Investigate tighter approximations (better abstract domains, better algorithm)
- Actually parallelize workflow across multiple cores/machines

1. De Nicola, Di Stefano, Inverso. Multi-agent systems with virtual stigmergy. *Sci. Comput. Program.* 187 (2020). DOI: <https://doi.org/10.1016/j.scico.2019.102345>
General introduction to LAbS
2. Di Stefano, De Nicola, Inverso. Verification of Distributed Systems via Sequential Emulation. *TOSEM* 31 (2022). DOI: <https://doi.org/10.1145/3490387>
Describes our general approach to LAbS verification
3. Di Stefano and Lang. Verifying Temporal Properties of Stigmergic Collective Systems Using CADP. In *ISoLA2021*. DOI: https://doi.org/10.1007/978-3-030-89159-6_29
Baseline CADP-based verification workflow and benchmark description

Backup slides

Evaluation (x)

- Read $val(x)$
- Mark x for a
qry-message

Evaluation (x)

- Read $val(x)$
- Mark x for a qry-message

Assignment ($x \leftarrow v$)

- Compute the current timestamp t
- Record $val(x) \leftarrow v, ts(x) \leftarrow t$
- Mark x for a put-message
- Unmark x for qry

Evaluation (x)

- Read $val(x)$
- Mark x for a qry-message

Assignment ($x \leftarrow v$)

- Compute the current timestamp t
- Record $val(x) \leftarrow v, ts(x) \leftarrow t$
- Mark x for a put-message
- Unmark x for qry

Messaging

- Messages are sent asynchronously to all neighbours
- Neighbourhood is defined as a predicate on sender and potential receiver
- Different variables may have different predicates

Receiving $\langle \text{put}, x, v, t \rangle$

- If $t > ts(x)$:
 1. Record $val(x) \leftarrow v, ts(x) \leftarrow t$
 2. Mark x for a put-message
- Otherwise, ignore the message

Receiving $\langle \text{qry}, x, v, t \rangle$

- Mark x for a put-message
- If $t > ts(x)$, then record $val(x) \leftarrow v, ts(x) \leftarrow t$

Table: Time and memory requirements for the *Compositional* and *Parallel* workflows.

	Compositional	Parallel
Time	$\sum_{Tasks} time(\mathcal{T})$	$\max_i \{time(\mathcal{T}_i)\} + time(\mathcal{T}_{\mathbb{P}}) + time(\mathcal{T}_{\mathbb{F}})$
Memory	$\max_{Tasks} mem(\mathcal{T})$	$\max \{ \sum_i mem(\mathcal{T}_i), mem(\mathcal{T}_{\mathbb{P}}), mem(\mathcal{T}_{\mathbb{F}}) \}$




Figure: Example of a diamond when 3 agents perform independent actions a_0 , a_1 , and a_2 . Dotted transitions are cut by applying the priority relation $a_0 \succ a_1 \succ a_2$.